

In [232...

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import norm
import seaborn as sns
from matplotlib.pyplot import figure
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
from sklearn.utils import resample
from sklearn.metrics import accuracy_score
```

In [233...

```
df = pd.read_csv("walmart_data.csv")
```

Problem Statement

- **Primary Goal**
 - Recognizing **Purchase pattern** of Products wrt. Gender , Age , Occupation , Marital_Status, City_Category etc.
 - Identifying **customer segments, profiling and formulating marketing strategy**
 - How to **drive sales of products and revenue** , across product categories
 - Data driven discounting / offers among customer segments
- **Statistical summary**
 - More **likelihood of purchase**
 - Range / Limitation of data
- **Long term benefits** : Sales growth , Customer acquisition and retention

In [234...

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   User_ID               550068 non-null  int64  
 1   Product_ID            550068 non-null  object  
 2   Gender                550068 non-null  object  
 3   Age                   550068 non-null  object  
 4   Occupation            550068 non-null  int64  
 5   City_Category         550068 non-null  object  
 6   Stay_In_Current_City_Years  550068 non-null  object  
 7   Marital_Status        550068 non-null  int64  
 8   Product_Category      550068 non-null  int64  
 9   Purchase              550068 non-null  int64  
dtypes: int64(5), object(5)
memory usage: 42.0+ MB
```

Basic Analysis

- **Analysing metrics** - Basic metrics
 - Observations on **shape** of data
 - **Data types** of all the attributes

- **Conversion** of categorical attributes to 'category' (If required)
- **Structure & characteristics** of the dataset
- Statistical summary

In [235...

df.shape

Out[235]: (550068, 10)

In [236...

df.describe()

Out[236]:

	User_ID	Occupation	Marital_Status	Product_Category	Purchase
count	5.500680e+05	550068.000000	550068.000000	550068.000000	550068.000000
mean	1.003029e+06	8.076707	0.409653	5.404270	9263.968713
std	1.727592e+03	6.522660	0.491770	3.936211	5023.065394
min	1.000001e+06	0.000000	0.000000	1.000000	12.000000
25%	1.001516e+06	2.000000	0.000000	1.000000	5823.000000
50%	1.003077e+06	7.000000	0.000000	5.000000	8047.000000
75%	1.004478e+06	14.000000	1.000000	8.000000	12054.000000
max	1.006040e+06	20.000000	1.000000	20.000000	23961.000000

In [237...

df.describe(include=object)

Out[237]:

	Product_ID	Gender	Age	City_Category	Stay_In_Current_City_Years
count	550068	550068	550068	550068	550068
unique	3631	2	7	3	5
top	P00265242	M	26-35	B	1
freq	1880	414259	219587	231173	193821

Non-Graphical Analysis

In [238...

df.head()

Out[238]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Proc
0	1000001	P00069042	F	0-17	10	A	2	0	
1	1000001	P00248942	F	0-17	10	A	2	0	
2	1000001	P00087842	F	0-17	10	A	2	0	
3	1000001	P00085442	F	0-17	10	A	2	0	
4	1000002	P00285442	M	55+	16	C	4+	0	

```
In [239... df["Product_ID"].value_counts()
```

```
Out[239]: P00265242    1880
          P00025442    1615
          P00110742    1612
          P00112142    1562
          P00057642    1470
          ...
          P00314842     1
          P00298842     1
          P00231642     1
          P00204442     1
          P00066342     1
          Name: Product_ID, Length: 3631, dtype: int64
```

```
In [240... df["User_ID"].value_counts()
```

```
Out[240]: 1001680    1026
          1004277     979
          1001941     898
          1001181     862
          1000889     823
          ...
          1002690       7
          1002111       7
          1005810       7
          1004991       7
          1000708       6
          Name: User_ID, Length: 5891, dtype: int64
```

```
In [241... df["Product_Category"].value_counts()
```

```
Out[241]: 5      150933
          1      140378
          8      113925
          11     24287
          2      23864
          6      20466
          3      20213
          4      11753
          16     9828
          15     6290
          13     5549
          10     5125
          12     3947
          7      3721
          18     3125
          20     2550
          19     1603
          14     1523
          17      578
          9       410
          Name: Product_Category, dtype: int64
```

```
In [242... df["Gender"].value_counts()
```

```
Out[242]: M      414259
          F      135809
          Name: Gender, dtype: int64
```

```
In [243... df["Marital_Status"].value_counts()
```

```
Out[243]: 0      324731
          1      225337
          Name: Marital_Status, dtype: int64
```

```
In [244]: df["Age"].value_counts() # different generation categories
```

```
Out[244]: 26-35      219587
          36-45      110013
          18-25       99660
          46-50       45701
          51-55       38501
          55+        21504
          0-17        15102
          Name: Age, dtype: int64
```

- **Data type conversion**

```
In [245]: df.User_ID = df.User_ID.astype(object)
          df.Occupation = df.Occupation.astype(object)
          df.Product_Category = df.Product_Category.astype(object)
```

```
In [246]: def marital_status_to_category(marital_status):
          if marital_status == 0:
              return "Un-Married"
          else:
              return "Married"
```

```
In [247]: df["Marital_Status"] = df["Marital_Status"].apply(marital_status_to_category)
```

```
In [248]: df["Marital_Status"].value_counts()
```

```
Out[248]: Un-Married      324731
          Married        225337
          Name: Marital_Status, dtype: int64
```

```
In [249]: df["Stay_In_Current_City_Years"].value_counts()
```

```
Out[249]: 1      193821
          2      101838
          3       95285
          4+      84726
          0       74398
          Name: Stay_In_Current_City_Years, dtype: int64
```

- **Feature types by data type**

```
In [250]: continious_features = df.select_dtypes(include=['int64', 'float64']).columns
          continious_features
```

```
Out[250]: Index(['Purchase'], dtype='object')
```

```
In [251]: categorical_features = df.select_dtypes(exclude=['int64', 'float64']).columns
          categorical_features
```

```
Out[251]: Index(['User_ID', 'Product_ID', 'Gender', 'Age', 'Occupation', 'City_Category',  
        'Stay_In_Current_City_Years', 'Marital_Status', 'Product_Category'],  
        dtype='object')
```

```
In [252... def percentage_wise_features(feature_list):  
    for feature_name in feature_list:  
        features_stats = df[feature_name].value_counts(normalize=True)*100  
        print("<<< ", feature_name, ">>> ")  
        print()  
        print(features_stats)  
        print()
```

```
In [253... percentage_wise_features(categorical_features)
```

```
<<< User_ID >>>

1001680    0.186522
1004277    0.177978
1001941    0.163253
1001181    0.156708
1000889    0.149618
...
1002690    0.001273
1002111    0.001273
1005810    0.001273
1004991    0.001273
1000708    0.001091
Name: User_ID, Length: 5891, dtype: float64
```

```
<<< Product_ID >>>

P00265242    0.341776
P00025442    0.293600
P00110742    0.293055
P00112142    0.283965
P00057642    0.267240
...
P00314842    0.000182
P00298842    0.000182
P00231642    0.000182
P00204442    0.000182
P00066342    0.000182
Name: Product_ID, Length: 3631, dtype: float64
```

```
<<< Gender >>>

M    75.310507
F    24.689493
Name: Gender, dtype: float64
```

```
<<< Age >>>

26-35    39.919974
36-45    19.999891
18-25    18.117760
46-50     8.308246
51-55     6.999316
55+       3.909335
0-17      2.745479
Name: Age, dtype: float64
```

```
<<< Occupation >>>

4    13.145284
0    12.659889
7    10.750125
1     8.621843
17    7.279645
20    6.101427
12    5.668208
14    4.964659
2     4.833584
16    4.612339
6     3.700452
3     3.208694
10    2.350618
5     2.213726
15    2.211545
11    2.106285
19    1.538173
```

```
13      1.404917
18      1.203851
9       1.143677
8       0.281056
Name: Occupation, dtype: float64
```

```
<<<  City_Category  >>>
```

```
B      42.026259
C      31.118880
A      26.854862
Name: City_Category, dtype: float64
```

```
<<<  Stay_In_Current_City_Years  >>>
```

```
1      35.235825
2      18.513711
3      17.322404
4+     15.402823
0      13.525237
Name: Stay_In_Current_City_Years, dtype: float64
```

```
<<<  Marital_Status  >>>
```

```
Un-Married    59.034701
Married       40.965299
Name: Marital_Status, dtype: float64
```

```
<<<  Product_Category  >>>
```

```
5      27.438971
1      25.520118
8      20.711076
11     4.415272
2      4.338373
6      3.720631
3      3.674637
4      2.136645
16     1.786688
15     1.143495
13     1.008784
10     0.931703
12     0.717548
7      0.676462
18     0.568112
20     0.463579
19     0.291419
14     0.276875
17     0.105078
9      0.074536
Name: Product_Category, dtype: float64
```

In [254...

```
categorical_features
```

Out[254]:

```
Index(['User_ID', 'Product_ID', 'Gender', 'Age', 'Occupation', 'City_Category',
      'Stay_In_Current_City_Years', 'Marital_Status', 'Product_Category'],
      dtype='object')
```

Multi-feature Analysis (Marginal Probability)

- **Gender based**

In [255...

```
pd.crosstab(df['Product_Category'], df['Gender'], margins=True, normalize=True)*100
```

Out[255]:

	Gender	F	M	All
Product_Category				
	1	4.514169	21.005948	25.520118
	2	1.028600	3.309773	4.338373
	3	1.091865	2.582772	3.674637
	4	0.661555	1.475090	2.136645
	5	7.628330	19.810642	27.438971
	6	0.828807	2.891824	3.720631
	7	0.171433	0.505028	0.676462
	8	6.100700	14.610375	20.711076
	9	0.012726	0.061811	0.074536
	10	0.211247	0.720456	0.931703
	11	0.861530	3.553742	4.415272
	12	0.278511	0.439037	0.717548
	13	0.265785	0.742999	1.008784
	14	0.113259	0.163616	0.276875
	15	0.190158	0.953337	1.143495
	16	0.436673	1.350015	1.786688
	17	0.011271	0.093807	0.105078
	18	0.069446	0.498666	0.568112
	19	0.081990	0.209429	0.291419
	20	0.131438	0.332141	0.463579
	All	24.689493	75.310507	100.000000

• Age based

In [256...

```
pd.crosstab(df['Product_Category'], df['Age'], margins=True, normalize=True)*100
```


Out[256]:

	Age	0-17	18-25	26-35	36-45	46-50	51-55	55+	All
Product_Category									
1		0.651738	4.901576	10.589418	5.026288	1.904128	1.645069	0.801901	25.520118
2		0.146346	0.804991	1.623072	0.892981	0.382680	0.323778	0.164525	4.338373
3		0.218155	0.856258	1.392919	0.700641	0.250151	0.167979	0.088535	3.674637
4		0.137801	0.447763	0.762088	0.427947	0.179978	0.123257	0.057811	2.136645
5		0.787175	5.185177	11.175527	5.340612	2.176276	1.798505	0.975698	27.438971
6		0.072536	0.681552	1.542537	0.708821	0.294873	0.263604	0.156708	3.720631
7		0.009635	0.087444	0.300145	0.147073	0.059447	0.048358	0.024361	0.676462
8		0.410495	3.256143	8.045551	4.235113	1.937215	1.697972	1.128588	20.711076
9		0.002909	0.011453	0.027997	0.019452	0.005999	0.005272	0.001454	0.074536
10		0.020179	0.109623	0.324869	0.224518	0.094534	0.094352	0.063628	0.931703
11		0.134529	0.835715	1.795051	0.900434	0.382498	0.265058	0.101987	4.415272
12		0.022724	0.079808	0.199248	0.180705	0.094534	0.078718	0.061811	0.717548
13		0.020361	0.137438	0.381044	0.227245	0.100169	0.087807	0.054721	1.008784
14		0.007090	0.041813	0.102533	0.056720	0.027088	0.027997	0.013635	0.276875
15		0.029087	0.186159	0.431219	0.253605	0.109441	0.092352	0.041631	1.143495
16		0.041631	0.290510	0.748635	0.355411	0.159798	0.122167	0.068537	1.786688
17		0.001091	0.007454	0.023088	0.024542	0.017271	0.019452	0.012180	0.105078
18		0.004908	0.061629	0.189431	0.127621	0.063810	0.076900	0.043813	0.568112
19		0.010726	0.049994	0.102351	0.058175	0.027088	0.024361	0.018725	0.291419
20		0.016362	0.085262	0.163253	0.091989	0.041268	0.036359	0.029087	0.463579
All		2.745479	18.117760	39.919974	19.999891	8.308246	6.999316	3.909335	100.000000

• Occupation based

In [257...

```
pd.crosstab(df['Product_Category'], df['Occupation'], margins=True, normalize=True)*100
```

Out[257]:

	Occupation	0	1	2	3	4	5	6	7	8	
	Product_Category										
	1	3.207422	1.875223	1.043871	0.717002	3.497386	0.662827	0.836442	2.915458	0.093079	0.000000
	2	0.540115	0.349411	0.201793	0.121076	0.552477	0.107078	0.149254	0.456489	0.017998	0.000000
	3	0.479032	0.283420	0.174160	0.111623	0.601017	0.081263	0.136347	0.291600	0.012180	0.000000
	4	0.269239	0.176887	0.093079	0.072355	0.309598	0.045813	0.080899	0.199793	0.007090	0.000000
	5	3.451391	2.394068	1.384011	0.960972	3.719358	0.595563	1.034599	2.839467	0.067446	0.000000
	6	0.466851	0.321233	0.188886	0.118349	0.483031	0.068901	0.133620	0.418676	0.008181	0.000000
	7	0.101624	0.067992	0.041631	0.024906	0.081626	0.006726	0.026542	0.059447	0.000909	0.000000
	8	2.560956	2.181185	1.139132	0.698641	2.514235	0.350684	0.869892	2.294262	0.047812	0.000000
	9	0.008908	0.004727	0.002182	0.002363	0.009272	0.001636	0.002909	0.007090	0.000182	0.000000
	10	0.115622	0.093625	0.051448	0.035450	0.086171	0.015453	0.034359	0.095443	0.002363	0.000000
	11	0.682097	0.305599	0.209247	0.125984	0.580292	0.143800	0.155435	0.438491	0.011453	0.000000
	12	0.080899	0.078899	0.044176	0.032541	0.067628	0.014180	0.029269	0.084535	0.000545	0.000000
	13	0.130893	0.100715	0.044904	0.039995	0.107441	0.017452	0.037268	0.110532	0.001454	0.000000
	14	0.035632	0.030723	0.014725	0.012362	0.032723	0.004181	0.013089	0.029087	0.000000	0.000000
	15	0.130529	0.096170	0.055266	0.031087	0.135256	0.026179	0.035632	0.114895	0.004363	0.000000
	16	0.222700	0.142528	0.081444	0.062174	0.216700	0.038722	0.075627	0.212337	0.002909	0.000000
	17	0.011453	0.011090	0.002545	0.002727	0.008363	0.001636	0.003272	0.015089	0.000182	0.000000
	18	0.077445	0.043267	0.027451	0.018361	0.048721	0.017271	0.016543	0.083808	0.001273	0.000000
	19	0.034178	0.025270	0.011271	0.008908	0.032360	0.004363	0.011453	0.034905	0.000909	0.000000
	20	0.052903	0.039813	0.022361	0.011817	0.061629	0.009999	0.017998	0.048721	0.000727	0.000000
	All	12.659889	8.621843	4.833584	3.208694	13.145284	2.213726	3.700452	10.750125	0.281056	1.000000

21 rows × 22 columns

• Marital_Status based

In [258...

```
pd.crosstab(df['Product_Category'], df['Marital_Status'], margins=True, normalize=True)*100
```

Out[258]:

	Marital_Status	Married	Un-Married	All
Product_Category				
1	10.181105	15.339013	25.520118	
2	1.768145	2.570228	4.338373	
3	1.427823	2.246813	3.674637	
4	0.831897	1.304748	2.136645	
5	11.139895	16.299076	27.438971	
6	1.513813	2.206818	3.720631	
7	0.305599	0.370863	0.676462	
8	8.819637	11.891439	20.711076	
9	0.029633	0.044904	0.074536	
10	0.426675	0.505028	0.931703	
11	1.748693	2.666579	4.415272	
12	0.347775	0.369772	0.717548	
13	0.433946	0.574838	1.008784	
14	0.123076	0.153799	0.276875	
15	0.484849	0.658646	1.143495	
16	0.748089	1.038599	1.786688	
17	0.050903	0.054175	0.105078	
18	0.269785	0.298327	0.568112	
19	0.119440	0.171979	0.291419	
20	0.194521	0.269058	0.463579	
All	40.965299	59.034701	100.000000	

- City_Category based

In [259...

```
pd.crosstab(df['Product_Category'], df['City_Category'], margins=True, normalize=True) * 100
```

Out[259]:

	City_Category	A	B	C	All
	Product_Category				
	1	6.377575	10.590145	8.552397	25.520118
	2	1.116407	1.898674	1.323291	4.338373
	3	0.898616	1.561080	1.214941	3.674637
	4	0.554477	0.950064	0.632104	2.136645
	5	7.673779	11.660013	8.105180	27.438971
	6	1.001149	1.549990	1.169492	3.720631
	7	0.222882	0.290691	0.162889	0.676462
	8	5.850004	8.644931	6.216141	20.711076
	9	0.019998	0.031632	0.022906	0.074536
	10	0.242334	0.375045	0.314325	0.931703
	11	1.200033	1.906128	1.309111	4.415272
	12	0.193249	0.304508	0.219791	0.717548
	13	0.293418	0.412858	0.302508	1.008784
	14	0.087444	0.114895	0.074536	0.276875
	15	0.312143	0.479577	0.351775	1.143495
	16	0.517754	0.734091	0.534843	1.786688
	17	0.021997	0.048539	0.034541	0.105078
	18	0.136892	0.252514	0.178705	0.568112
	19	0.049630	0.083990	0.157799	0.291419
	20	0.085080	0.136892	0.241606	0.463579
	All	26.854862	42.026259	31.118880	100.000000

- Stay_In_Current_City_Years based

In [260...

```
pd.crosstab(df['Product_Category'], df['Stay_In_Current_City_Years'], margins=True,norma
```

Out[260]: Stay_In_Current_City_Years 0 1 2 3 4+ All

Product_Category							
1	3.373219	8.875085	4.819950	4.541439	3.910426	25.520118	
2	0.581746	1.523266	0.838078	0.768996	0.626286	4.338373	
3	0.507028	1.274024	0.711185	0.651738	0.530662	3.674637	
4	0.288510	0.745181	0.387952	0.382135	0.332868	2.136645	
5	3.744992	9.653716	5.109550	4.794316	4.136398	27.438971	
6	0.497575	1.320928	0.678643	0.625559	0.597926	3.720631	
7	0.101442	0.234698	0.109259	0.125439	0.105623	0.676462	
8	2.788928	7.481620	3.679181	3.467571	3.293775	20.711076	
9	0.011271	0.023997	0.015816	0.013998	0.009453	0.074536	
10	0.120894	0.336867	0.165071	0.165979	0.142891	0.931703	
11	0.633558	1.480726	0.837351	0.727910	0.735727	4.415272	
12	0.097806	0.255968	0.133256	0.123439	0.107078	0.717548	
13	0.140892	0.364682	0.177978	0.171433	0.153799	1.008784	
14	0.039268	0.104533	0.045267	0.042904	0.044904	0.276875	
15	0.166525	0.398133	0.210156	0.191249	0.177433	1.143495	
16	0.239970	0.641739	0.327959	0.295236	0.281783	1.786688	
17	0.015634	0.030542	0.020907	0.020907	0.017089	0.105078	
18	0.077445	0.219609	0.095625	0.088898	0.086535	0.568112	
19	0.036177	0.110350	0.053812	0.046176	0.044904	0.291419	
20	0.062356	0.160162	0.096715	0.077081	0.067264	0.463579	
All	13.525237	35.235825	18.513711	17.322404	15.402823	100.000000	

Multi-feature Analysis (Conditional Probability)

```
In [261]: pd.crosstab(index=df['Gender'], columns=df['Product_Category'], margins=True, normalize=True)
```

Product_Category										
Gender										
F	0.182838	0.041661	0.044224	0.026795	0.308971	0.033569	0.006944	0.247097	0.000515	0.0085
M	0.278925	0.043948	0.034295	0.019587	0.263053	0.038399	0.006706	0.194002	0.000821	0.0095
All	0.255201	0.043384	0.036746	0.021366	0.274390	0.037206	0.006765	0.207111	0.000745	0.0093

```
In [262]: pd.crosstab(index=df['Age'], columns=df['Product_Category'], margins=True, normalize=True)
```

Out[262]:

Product_Category	1	2	3	4	5	6	7	8	9		
Age											
0-17	0.237386	0.053304	0.079460	0.050192	0.286717	0.026420	0.003509	0.149517	0.001059	0.0073	
18-25	0.270540	0.044431	0.047261	0.024714	0.286193	0.037618	0.004826	0.179721	0.000632	0.0060	
26-35	0.265266	0.040658	0.034893	0.019090	0.279948	0.038641	0.007519	0.201542	0.000701	0.0081	
36-45	0.251316	0.044649	0.035032	0.021397	0.267032	0.035441	0.007354	0.211757	0.000973	0.0112	
46-50	0.229185	0.046060	0.030109	0.021663	0.261942	0.035492	0.007155	0.233168	0.000722	0.0113	
51-55	0.235033	0.046259	0.023999	0.017610	0.256954	0.037661	0.006909	0.242591	0.000753	0.0134	
55+	0.205125	0.042085	0.022647	0.014788	0.249581	0.040086	0.006231	0.288690	0.000372	0.0162	
All	0.255201	0.043384	0.036746	0.021366	0.274390	0.037206	0.006765	0.207111	0.000745	0.0093	

In [263...

```
pd.crosstab(index=df['Marital_Status'], columns=df['Product_Category'], margins=True, no
```

Out[263]:

Product_Category	1	2	3	4	5	6	7	8	9	
Marital_Status										
Married	0.248530	0.043162	0.034854	0.020307	0.271935	0.036954	0.007460	0.215295	0.000723	0.0104
Un-Married	0.259830	0.043538	0.038059	0.022101	0.276093	0.037382	0.006282	0.201431	0.000761	0.0085
All	0.255201	0.043384	0.036746	0.021366	0.274390	0.037206	0.006765	0.207111	0.000745	0.0093

In [264...

```
pd.crosstab(index=df['City_Category'], columns=df['Product_Category'], margins=True, no
```

Out[264]:

Product_Category	1	2	3	4	5	6	7	8	9	
City_Category										
A	0.237483	0.041572	0.033462	0.020647	0.285750	0.037280	0.008299	0.217838	0.000745	0.0090
B	0.251989	0.045178	0.037145	0.022606	0.277446	0.036881	0.006917	0.205703	0.000753	0.0089
C	0.274830	0.042524	0.039042	0.020313	0.260459	0.037581	0.005234	0.199755	0.000736	0.0101
All	0.255201	0.043384	0.036746	0.021366	0.274390	0.037206	0.006765	0.207111	0.000745	0.0093

In [265...

```
pd.crosstab(index=df['Stay_In_Current_City_Years'], columns=df['Product_Category'], mar
```

Out[265]:

Product_Category	1	2	3	4	5	6	7	8	9	
Stay_In_Current_City_Years										
0	0.249402	0.043012	0.037488	0.021331	0.276889	0.036789	0.007500	0.206202	0.000833	
1	0.251877	0.043231	0.036157	0.021148	0.273974	0.037488	0.006661	0.212330	0.000681	
2	0.260345	0.045268	0.038414	0.020955	0.275987	0.036656	0.005902	0.198727	0.000854	
3	0.262171	0.044393	0.037624	0.022060	0.276770	0.036113	0.007241	0.200178	0.000808	
4+	0.253877	0.040660	0.034452	0.021611	0.268548	0.038819	0.006857	0.213842	0.000614	
All	0.255201	0.043384	0.036746	0.021366	0.274390	0.037206	0.006765	0.207111	0.000745	

- **Observation from Non-graphical Analysis**
 - **Gender - 1, 5, 8 product categories** are more likely to be purchased by **males than females**
 - **Age - 1,5 product categories** are more **frequently** being **purchased** in **age group 26-35**
 - **Marital_Status - 1, 5, 8 product categories** are more popular purchase among **un-married** users
 - **City Category**
 - **City A : 1, 5, 8 product categories** are popular
 - **City B : 1, 5, 8 product categories** are popular
 - **City C : 1, 5, 8 product categories** are popular
 - **Stay_In_Current_City_Years - 1, 5, 8 product categories** are popular

Missing Value & Outlier Detection

- **Missing value detection**

In [266...

```
percent_missing = df.isnull().sum() * 100 / len(df)
missing_value_df = pd.DataFrame({'column_name': df.columns,
                                'percent_missing': percent_missing})
missing_value_df.sort_values('percent_missing', ascending=False)
```

Out[266]:

	column_name	percent_missing
User_ID	User_ID	0.0
Product_ID	Product_ID	0.0
Gender	Gender	0.0
Age	Age	0.0
Occupation	Occupation	0.0
City_Category	City_Category	0.0
Stay_In_Current_City_Years	Stay_In_Current_City_Years	0.0
Marital_Status	Marital_Status	0.0
Product_Category	Product_Category	0.0
Purchase	Purchase	0.0

- **Outlier detection**

In [267...

```
def find_outliers_IQR(column_name):
    print("Outliers by feature name --> ", column_name)
    Q1=df[column_name].quantile(0.25)
    Q3=df[column_name].quantile(0.75)

    IQR=Q3-Q1
    lower = Q1 - 1.5*IQR
    upper = Q3 + 1.5*IQR

    outliers = df[((df[column_name]<lower) | (df[column_name]>upper))]

    return outliers
```

In [268...

```
outlier_df_by_purchase = find_outliers_IQR("Purchase")
outlier_df_by_purchase
```

Outliers by feature name --> Purchase

Out[268]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
343	1000058	P00117642	M	26-35	2	B	3	Un-Married
375	1000062	P00119342	F	36-45	3	A	1	Un-Married
652	1000126	P00087042	M	18-25	9	B	1	Un-Married
736	1000139	P00159542	F	26-35	20	C	2	Un-Married
1041	1000175	P00052842	F	26-35	2	B	1	Un-Married
...
544488	1005815	P00116142	M	26-35	20	B	1	Un-Married
544704	1005847	P00085342	F	18-25	4	B	2	Un-Married
544743	1005852	P00202242	F	26-35	1	A	0	Married
545663	1006002	P00116142	M	51-55	0	C	1	Married
545787	1006018	P00052842	M	36-45	1	C	3	Un-Married

2677 rows × 10 columns

In [269...]

```
outlier_df_by_occupation = find_outliers_IQR("Occupation")
outlier_df_by_occupation
```

Outliers by feature name --> Occupation

Out[269]:

User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Produ
---------	------------	--------	-----	------------	---------------	----------------------------	----------------	-------

In [270...]

```
outlier_df_by_Product_Category = find_outliers_IQR("Product_Category")
outlier_df_by_Product_Category
```

Outliers by feature name --> Product_Category

Out[270]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
545915	1000001	P00375436	F	0-17	10	A	2	Un-Married
545916	1000002	P00372445	M	55+	16	C	4+	Un-Married
545917	1000004	P00375436	M	46-50	7	B	2	Married
545918	1000006	P00375436	F	51-55	9	A	1	Un-Married
545919	1000007	P00372445	M	36-45	1	B	1	Married
...
550063	1006033	P00372445	M	51-55	13	B	1	Married
550064	1006035	P00375436	F	26-35	1	C	3	Un-Married
550065	1006036	P00375436	F	26-35	15	B	4+	Married
550066	1006038	P00375436	F	55+	1	C	2	Un-Married
550067	1006039	P00371644	F	46-50	0	B	4+	Married

4153 rows × 10 columns


Visual Analysis

In [271...

```
sns.kdeplot(data=df, x="Purchase")
```


Out[271]:

```
<AxesSubplot: xlabel='Purchase', ylabel='Density'>
```



In [272...

```
sns.set(font_scale = 1.1)
plt.figure(figsize=(10,5))
plt.xlabel("Gender", fontsize=17)
plt.ylabel("Density", fontsize=17)
sns.distplot(df["Product_Category"])
plt.title("Product_Category based distribution", fontdict = {"fontsize": 17})
plt.show()
```







Purchase Analysis


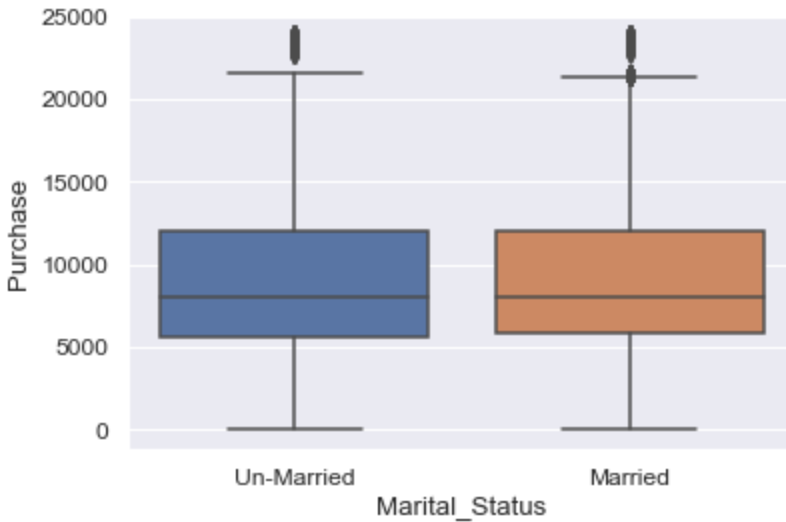
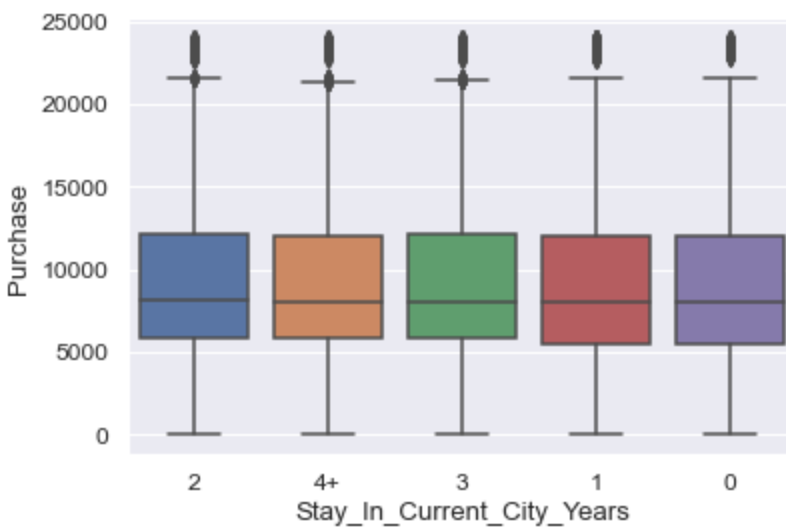
In [273... `selected_features = ['Gender', 'Age', 'Occupation', 'City_Category', 'Stay_In_Current_City_Years']`

In [274... `def boxplot_all_catogorical_vs_all_continious_features(feature_y_catg_list, feature_y_con):
 plt.figure()
 for categorical_feature in feature_y_catg_list:
 for continious_feature in feature_y_contn_list:
 sns.boxplot(x = categorical_feature, y=continious_feature, data=df)
 plt.show()`

In [275... `boxplot_all_catogorical_vs_all_continious_features(selected_features, continious_features)`







- **Observation from Visual Analysis**

- Except in "City C", **overall purchase pattern almost similar across 'Gender', 'Age', 'Stay_In_Current_City_Years', 'Marital_Status'**
- There is **varied purchase pattern** across **'Occupation', 'Product_Category'**
- **Total purchase for product category 6,7,9,10,15,16 are higher** than other product categories


- **Pair plot**

In [276...

```
selected_features = ['Gender', 'Age', 'Occupation', 'City_Category', 'Stay_In_Current_City_Years', 'Product_Category', 'Purchase']
```

```
In [277... sns.pairplot(data=df[selected_features])
plt.plot()
```

```
Out[277]: []
```



```
In [278... df_product_category_1_5_8 = df[(df["Product_Category"] == 1) | (df["Product_Category"] == 5) | (df["Product_Category"] == 8)]
df_product_category_1 = df[df["Product_Category"] == 1]
df_product_category_5 = df[df["Product_Category"] == 5]
df_product_category_8 = df[df["Product_Category"] == 8]
```

```
In [279... df_product_category_1_5_8["Product_Category"].value_counts()
```

```
Out[279]: 5    150933
1    140378
8    113925
Name: Product_Category, dtype: int64
```

```
In [280... df_product_category_1 = df[df["Product_Category"] == 1]
df_product_category_5 = df[df["Product_Category"] == 5]
df_product_category_8 = df[df["Product_Category"] == 8]
```

Product_Category Analysis

```
In [281... df_selected = df_product_category_1_5_8[selected_features]
```

```
In [282... def show_values_on_bars(axes, h_v="v", space=1):
    total = float(len(df_selected))
    def _show_on_single_plot(ax):
        if h_v == "v":
            for p in ax.patches:
                _x = p.get_x() + p.get_width() / 2
                _y = p.get_y() + p.get_height()
                value='{: .1f}%'.format(100 * p.get_height() / total)
                ax.text(_x, _y, value, ha="center")
```

```

elif h_v == "h":
    for p in ax.patches:
        _x = p.get_x() + p.get_width() + float(space)
        _y = p.get_y() + p.get_height()
        value='{:.1f}%'.format(100 * p.get_width()/total)
        ax.text(_x, _y, value, ha="left")
if isinstance(axs, np.ndarray):
    for idx, ax in np.ndenumerate(axs):
        _show_on_single_plot(ax)
else:
    _show_on_single_plot(axs)

```

In [283...

```

def plot_products_by_feature(feature_list):
    for feature_name in feature_list:
        sns.set(style="whitegrid")
        sns.set(font_scale = 1.1)
        plt.figure(figsize=(15,5))
        ax = sns.countplot(x=feature_name,data=df_selected,hue="Product_Category",order=
        plt.xlabel(feature_name, fontsize=17)
        plt.ylabel("Product_Category",fontsize=17)
        plt.title(feature_name + " distribution",fontdict ={"fontsize": 17})
        show_values_on_bars(ax,h_v="v",space=1)



```

In [284...


```

plot_products_by_feature(['Gender', 'Age', 'Occupation', 'City_Category', 'Stay_In_Current

```




Occupation distribution

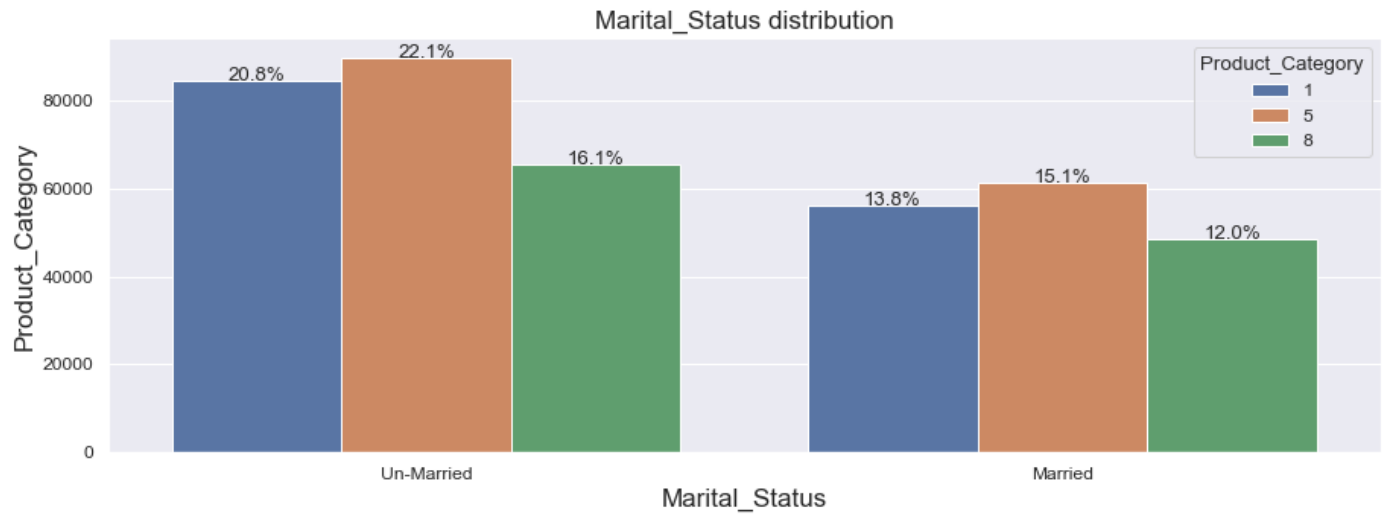


City_Category distribution



Stay_In_Current_City_Years distribution





- **Outlier treatment [Removing outliers]** just to ensure sample mean distribution is not impacted by outlier while using CLT]

In [285...

```
df.drop(outlier_df_by_purchase.index,inplace=True)  
df.drop(outlier_df_by_Product_Category.index,inplace=True)
```

- **Re-checking outliers post Outlier removal**

In [286...

```
outlier_df_by_purchase = find_outliers_IQR("Purchase")  
outlier_df_by_purchase
```

Outliers by feature name --> Purchase

Out[286]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status
5846	1000949	P00111042	M	51-55	17	B	4+	Married
10481	1001611	P00111742	M	26-35	0	B	1	Un-Married
11148	1001682	P00111742	M	26-35	2	A	3	Married
12497	1001884	P00111742	M	46-50	20	B	1	Married
14484	1002147	P00174242	F	18-25	0	B	3	Un-Married
...
534978	1004354	P00111742	M	36-45	14	B	2	Un-Married
541307	1005359	P00174242	M	18-25	16	B	3	Un-Married
542825	1005580	P00071442	M	46-50	7	B	4+	Un-Married
545664	1006002	P00071442	M	51-55	0	C	1	Married
545864	1006036	P00111042	F	26-35	15	B	4+	Married

165 rows × 10 columns

In [287...

```
outlier_df_by_occupation = find_outliers_IQR("Occupation")
outlier_df_by_occupation
```

Outliers by feature name --> Occupation

Out[287]:

User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Produ
---------	------------	--------	-----	------------	---------------	----------------------------	----------------	-------

In [288...

```
outlier_df_by_Product_Category = find_outliers_IQR("Product_Category")
outlier_df_by_Product_Category
```

Outliers by feature name --> Product_Category

Out[288]:

User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Produ
---------	------------	--------	-----	------------	---------------	----------------------------	----------------	-------

Confidence Interval (CI)

In [295...

```
def getConfidenceIntervalByFeatureValue(feature_x_name,feature_x_value,feature_y,bootst
    if dataset_catagory == "Product_Category_All":
        data = df
    else:
        data = df_product_category_1_5_8

    # Configure bootstrap
    transactions = data[data[feature_x_name] == feature_x_value][feature_y]

    sample_size = 10000
    bootstrap_size = int(bootstrap_repetition_factor *len(transactions))
```

[illegible]

In [296...

```
def getConfidenceIntervalByFeature(feature_x, feature_y, bootstrap_repetition_factor, dataset_category):
    if dataset_category == "Product_Category_All":
        data = df
    else:
        data = df_product_category_1_5_8



    levels = data[feature_x].unique()
    for level in levels:
        getConfidenceIntervalByFeatureValue(feature_x, level, feature_y, bootstrap_repetition_factor)
```

Gender (Confidence Interval)



In [297...

```
getConfidenceIntervalByFeature("Gender", "Purchase", 0.7, "Product Category All")
```

```
>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>
CI by Gender == F Mean   8745.09616804928    S.D -   15.159546487015946
```









```
>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>
CI by Marital Status == Un-Married Mean    9269.404120659836   S.D -   10.223493395464192
```




In [299...

[illegible]







```
C.I (on the mean) [9154.316921490956, 9228.21577729623]
```

C.I (on the percentile) by [9154.67105908 9227.92701592]

```
C.I length 73.89885580527334
```

[illegible]


CI for selected Product Category

In [300...

```
getConfidenceIntervalByFeature("Gender", "Purchase", 0.7, "Product_category_1_5_8")
```

[illegible]

CI by Gender == F Mean 8509.97255433127 S.D - 15.166278523648097




C.I (on the mean) [8480.24664842492, 8539.698460237621]

C.I (on the percentile) by [8480.22798135 8539.90933269]

C.I length 59.451811812701635


[illegible][illegible]

CI by Gender == M Mean 9354.916858801425 S.D - 9.740679785440566




In [301...

```
>>> CI by Marital Status == Un-Married Mean   9163.403381226546    S.D -   10.90388720319212
```





```
>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>>
CI by Marital Status == Married Mean   9119.884802769377 S.D.- 12.90513721186992
```








[illegible]

```
getConfidenceIntervalByFeature("Age", "Purchase", 0.7, "Product_category_1_5_8")
```



CI by Age == 26-35 Mean 9092.811687781397 S.D - 13.106393296943022





```
getConfidenceIntervalByFeature("City_Category", "Purchase", 0.7, "Product_category_1_5_8")
```

[illegible][illegible]

A histogram showing the distribution of years since 1980 for the 1000th birth cohort. The x-axis is labeled 'years since 1980' and ranges from 880 to 9060 with major ticks every 200 units. The y-axis represents frequency, ranging from 0 to 300 with major ticks every 50 units. The distribution is unimodal and roughly symmetric, centered around 9200 years since 1980. The peak frequency is approximately 340.

[illegible]

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 543238 entries, 0 to 545914
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               543238 non-null  object
1   Product_ID                            543238 non-null  object
2   Gender                                543238 non-null  object
3   Age                                    543238 non-null  object
4   Occupation                            543238 non-null  object
5   City_Category                         543238 non-null  object
6   Stay_In_Current_City_Years           543238 non-null  object
7   Marital_Status                        543238 non-null  object
8   Product_Category                     543238 non-null  object
9   Purchase                             543238 non-null  int64
dtypes: int64(1), object(9)
memory usage: 45.6+ MB
```

- Answering questions

- Are women spending more money per transaction than men? Why or Why not?
 - For **product category 1,3,5 - Mens are spending more money than women**
- Confidence intervals(CI) and distribution of the mean of the expenses by female and male customers
 - **Men** : CI ==> **[9416.093496109974, 9452.347891575842] length 36.25**
 - **Female** : CI ==> **[8715.38345693473, 8774.808879163831] length 59.42**
 - **Conclusion: More accurate estimation of expenses**(i.e. of population) **of Male** is possible **than that of Female** , as CI length is smaller(i.e 36.25) for men.
- Are confidence intervals of average male and female spending overlapping? How can Walmart leverage this conclusion to make changes or improvements?
 - **Confidence intervals** of average expenses by male and female customers spending **not overlapping**
 - **Conclusion** - The **null hypothesis of zero difference** between **expenses by male and female customers** will be **rejected**. Hence **there is difference in expenses by male and female customers**
- Results when the same activity is performed for Married vs Unmarried
 - **Married** : CI ==> **[9232.611558739225, 9280.799169773572] length 48.187611034347356**
 - **UnMarried**: CI ==> **[9249.366073604726, 9289.442167714946] length 40.076094110219856**
 - **Conclusion: Slitely better estimation of expenses** of **Un-Married people** than that of Married people as CI length is smaller for Un-Married people (i.e 40.07)
 - **Confidence intervals** of average expenses by Married and Unmarried customers spending **overlapping**
 - **Conclusion** - The **null hypothesis of zero difference** between **expenses by Married and Unmarried customers** will be **not be rejected**. Hence there is **no significant evidence** that there **ZERO difference in expenses by Married and Unmarried customers**
- Results when the same activity is performed for Age
 - **0-17** CI ==> [8859.372659501589, 9048.274073326213] length 188.9014138246239
 - **18-25** CI ==> [9154.316921490956, 9228.21577729623] length 73.89885580527334
 - **26-35** CI ==> [9228.872515157269, 9278.321016643575] length 49.44850148630576
 - **36-45** CI ==> [9287.944471233906, 9357.449021916815] length 69.5045506829083
 - **46-50** CI ==> [9150.26308149336, 9256.016655633373] length 105.75357414001337
 - **51-55** CI ==> [9444.769369355721, 9562.869448463816] length 118.10007910809509
 - **55+** CI ==> [9251.184228797556, 9406.240591906604] length 155.0563631090481

- **Conclusion** : Better estimation of expenses of population based on Age group **26-35** (length 49.45) , **36-45** (length 69.50), **18-25**(length 73.89)
 - **Confidence intervals** of average expenses by different age group customers spending **overlapping** , except group **0-17**
 - **Conclusion** - The **null hypothesis of zero difference** between **expenses by different age group customers** will be **not be rejected**. Hence there is **no significant evidence** that there **ZERO difference in expenses by different age group customers**
- **Final Insights**
 - **Insights based on exploration**
 - **Gender** based - **Males spend more** on **Product categories - 1, 5 and 8** than females
 - **Age** based - **Age group 26-35** spend more on **Product categories - 1, 5**
 - **Marital_Status** based - **1, 5 and 8 product categories** are more popular purchase among **un-married** customers
 - **City Category**
 - **Product categories - 1, 5 and 8** are **equally being purchased** in city category **A, B and C**
 - Except in "City C" , **overall purchase pattern almost similar across 'Gender', 'Age', 'Stay_In_Current_City_Years', 'Marital_Status'**
 - There is **varied purchase pattern** across **'Occupation','Product_Category'**
 - **Total purchase** for **product category 6,7,9,10,15,16** are **higher** than other product categories
 - **Insights based on CLT and generalization on Population**
 - **More accurate estimation of expenses**(i.e. of population) **of Male** is possible **than that of Female** , as CI length is smaller(i.e 36.25) for men.
 - There is a **definite difference in expenses by male and female customers**
 - **Slitely better estimation of expenses** possible of **Un-Married customers** than that of Married customers as CI length is smaller for Un-Married people (i.e 40.07)
 - There is **no significant evidence** that there **ZERO difference in expenses by Married and Unmarried customers**
 - Better estimation of expenses of population possible based on Age group **26-35** (length 49.45) , **36-45** (length 69.50), **18-25**(length 73.89)
 - **Comments on the distribution of the variables and relationship between them**
 - Product category 1,5,8 have highest distribution of expenses
 - The distribution seems to follow normal distribution
 - **Comments for each univariate and bivariate plots**
 - **Analysis based on top product categories i.e. 1, 5 and 8**
 - **Gender:**
 - **Females spend more** on product category **5,8**
 - **Males spend more** on product category **1,5**
 - **Age group:**
 - **78% purchases (of product category 1,5,8)** are contributed by age group **26-35(40.5%)** , **36-45 (19.7%)**, **18-25(18.1%)**
 - **26-35 spend more** on product category **5**
 - **55+ spend more** on product category **8**
 - **Occupation:**
 - **36.7% purchases (of product category 1,5,8)** are contributed across **occupation category 4, 0 and 7**

- **45.4% purchases (of product category 1,5,8)** are contributed across **occupation category 4, 0 , 7 and 1**
 - **City Category:**
 - **41.9% purchases (of product category 1,5,8)** are contributed by **city B**
 - **Product category 1,5,8** are more popular in city **B and C** , i.e. ~ **73 % of purchases of product 1,5,8**
 - **Product category 5 is more popular in city B, A**
 - **Product category 8 is more popular in city B**
 - **Product category 1 is more popular in city B and C**
 - **Stay_In_Current_City_Years:**
 - **35.3% purchases for product category 1,5,8 are contributed** by people **staying in the city for 1 year**
 - **Marital_Status:**
 - **Un-married** customers contributing to more purchase of **Product 1**
- **Recommendations**
- **Actionable items for business**
 - **Gender based targeted channel** would be effective to boost sales
 - **Males spend more** on product category **1,5**
 - More accurate positioning for Males
 - For **Females more focus on product category 5,8**
 - **Marital status** based targeted channels are very effective too
 - For **Un-married**, target **Product category 5**
 - Both **Un-married** and **married** clients, **target product category 1,8**
 - **Age group based targeted channels** can contribute to productive sales
 - For age group **26-35**, target product category **5**
 - More accurate positioning for 26-35 age group
 - For age group **36-45** and **18-25**, target for product category **1, 5 and 8**
 - **City category with Product category** based targeted channels will yield more results . **More focus on city category "B"** to multiple sales
 - Target **Product category 5 in city B, A**
 - Target **Product category 8 in city B**
 - Target **Product category 1 in city B and C**
 - Target **Product category 1,5,8** in city **B and C** [**~ 73 % of historical purchases of product 1,5,8**]